PLAYBOOK NO. 08 · AI AGENTS

Agents at work. Letting AI act, safely.

AI that does not just answer but gets things done. Where agents earn their place today, and the guardrails that keep them on track.

INSIDE

Contents

What AI agents are, where they already do useful work, and the permissions, approvals and checks that let them act on your behalf without nasty surprises.

Foreword 03

Introduction 04

01 What agents are, and what they can do 06

02 Choosing your first agents 11

03 Running agents safely 16

FOREWORD

The step from answering to acting is the biggest change in AI since the chatbot.

For most organisations, AI so far has meant an assistant that answers questions and drafts text. A person reads the output and decides what to do with it. Agents change that. They can look things up, update records, raise tickets, write and run code and work through a task from start to finish.

That is a real opportunity. Much of the work that slows teams down is not hard thinking but the steps around it: gathering information, copying it between systems, chasing updates. Agents are well suited to exactly that kind of work.

It also changes the risk. An assistant that gets something wrong produces a bad draft. An agent that gets something wrong can send the email, change the record or delete the file. This playbook shows how to get the benefit while keeping that risk small, visible and under your control.



The Transputec team
Transputec

INTRODUCTION

Give AI a job, not just a question.

An agent is AI with tools, a goal and permission to act. Done well, it takes routine work off your people's plates. Done carelessly, it creates a new kind of operational risk. The difference is almost entirely in how you set it up.

Start with narrow, repeatable work

The best first agents do one well-defined job with clear inputs, clear outputs and a person who already owns the process.

Set the level of autonomy deliberately

Decide whether the agent suggests, acts with approval or acts within limits. Start low and earn the right to go higher.

Treat agents like new staff with system access

Give them only the permissions they need, log what they do and review their work until they have proved themselves.

An agent should never be able to do more damage than the person who approved it.

Least privilege, human approval at the moments that matter and a full audit trail. Those three habits make agents safe enough to be useful.

01

PART 01

From answers to actions.

What makes something an agent, what agents are good at today, and how much freedom to give them.

01 Anatomy of an agent

02 What they do well today

03 Levels of autonomy

UNDERSTANDING AGENTS

Model, tools, goal, loop: the four parts of every agent

> **An assistant tells you what to do. An agent does it. That is why the setup matters so much more.**

Strip away the marketing and every AI agent has four parts. There is a **model**, the language model that reads, reasons and decides what to do next. There are **tools**, the connections that let it search, read a file, query a database, call an API or update a system. There is a **goal**, the outcome you have asked for. And there is a **loop**: the agent takes a step, looks at the result and decides on the next step until the goal is met or it gets stuck.

That loop is what separates an agent from an assistant. Ask an assistant to find overdue invoices and it will explain how. Ask an agent and it will query the finance system, compile the list, draft the reminders and, if permitted, send them.

The tools are where the value and the risk both live. An agent with read-only access to a knowledge base is low risk. An agent that can issue refunds, change user accounts or push code to production needs the same care you would apply to a new employee with the same access.

Agents are not magic. They can misread instructions, choose the wrong tool, loop without progress or report success when the job is only half done. They work best on tasks where the steps are fairly predictable, the result can be checked and a mistake can be caught and reversed.

Agents also come in different shapes. Some are built into products you already use, such as service desk, CRM or coding tools, and switched on with a setting. Others are assembled on an AI platform for one specific process. A growing number of general assistants can act as agents too, carrying out multi-step tasks in a browser or in connected apps. The shape affects who builds it, who supports it and where its controls live.

The practical question is therefore not whether to use agents, but which jobs, with which tools, at what level of autonomy. The rest of this playbook is about answering that question with confidence. Start narrow, keep a person accountable and widen the agent's freedom only when its track record justifies it.

THE AGENT LOOP

How an agent works through a task

Every agent, however sophisticated, repeats a version of this cycle until the job is done.

Understand the goal

The agent reads the request, the instructions it has been given and any context such as documents or ticket history.

Plan the next step

It decides what to do next: search, read, calculate, update a record or ask a person for help.

Use a tool

It calls one of the tools it has been granted, within the permissions that tool allows.

Check the result

It reads what came back and judges whether it is closer to the goal or needs another approach.

Finish or hand over

It completes the task and reports what it did, or stops and passes the work to a person with a summary.

> **Good agents know when to stop. Build in clear hand-over points so an agent asks rather than guesses.**

WHERE AGENTS HELP TODAY

Six jobs agents already do well

Agents are strongest on multi-step work that follows a recognisable pattern and ends in something a person can check.

Research and triage

Gathering information from several sources, sorting incoming requests and routing them to the right team with a summary.

Data entry and updates

Reading emails, forms or documents and updating the CRM, finance or HR system, with exceptions flagged for review.

Coding and IT operations

Writing and testing code changes, investigating alerts, gathering diagnostics and running approved fixes from a runbook.

Customer service workflows

Looking up orders, checking policy, drafting replies and completing routine actions such as rebooking or resending documents.

> If a capable new starter could do the task from a written procedure, an agent may be able to do it too, with the same supervision.

LEVELS OF AUTONOMY

How much freedom should an agent have?

Autonomy is a setting you choose, task by task. Most organisations should start at the top of this table and move down only with evidence.

LEVEL	WHAT THE AGENT DOES	THE PERSON'S ROLE	GOOD FOR
Inform	Gathers and summarises information. Changes nothing.	Reads and decides.	Research, monitoring, briefings.
Suggest	Proposes an action or prepares a draft.	Reviews, edits and carries it out.	Replies, code changes, record updates.
Act with approval	Prepares the action and waits for a yes.	Approves or rejects each action.	Refunds, account changes, outbound emails.
Act within limits	Acts alone inside set boundaries and reports back.	Sets limits and reviews logs.	Routine, reversible, high-volume tasks.
Act freely	Acts across systems with little oversight.	Rarely appropriate today.	Tightly sandboxed experiments only.

02

PART 02

Your first agents.

Where agents are paying their way, and how to pick a first use case that builds confidence rather than risk.

01 Use cases by function

02 Assistant or agent?

03 Picking the starting point

USE CASES

Research, triage and data entry

SALES AND BIDS

Account and tender research

THE TASK

Pulling together background on a prospect or a tender before a bid decision.

AUTONOMY

Inform. It changes nothing and a person makes the call.

WHAT THE AGENT DOES

Searches approved sources, reads the documents and compiles a structured brief.

WATCH OUT

Web content can be wrong or manipulated. Ask for sources and check key facts.

SHARED INBOXES

Triage and routing

THE TASK

Sorting a busy shared mailbox or request queue.

AUTONOMY

Act within limits. It can label and route, but not reply or delete.

WHAT THE AGENT DOES

Classifies each request, adds a summary and routes it to the right team or queue.

WATCH OUT

Sample its decisions weekly and keep an easy route for misrouted items.

FINANCE AND OPERATIONS

Data entry from documents

THE TASK

Keying invoices, orders or forms into business systems.

AUTONOMY

Act with approval at first, then within limits for clean matches.

WHAT THE AGENT DOES

Extracts the details, matches them to records and prepares the entry.

WATCH OUT

Anything that moves money stays behind an approval step.

USE CASES

Coding, IT operations and customer service

SOFTWARE TEAMS

Coding agents

THE TASK

Fixing bugs, writing tests and making routine code changes.

AUTONOMY

Suggest. A developer reviews and merges every change.

WHAT THE AGENT DOES

Reads the codebase, makes the change, runs the tests and opens a pull request.

WATCH OUT

No direct access to production or secrets. Normal code review still applies.

IT OPERATIONS

Alert investigation and runbooks

THE TASK

Working out why an alert fired and what to do about it.

AUTONOMY

Act with approval for fixes. Inform for diagnosis.

WHAT THE AGENT DOES

Gathers logs and metrics, checks recent changes and proposes a fix from the runbook.

WATCH OUT

Limit it to named, reversible runbook actions and log every one.

CUSTOMER SERVICE

End-to-end routine requests

THE TASK

Handling routine requests such as order status, rebooking or document resends.

AUTONOMY

Act within limits, with clear rules for when to hand over to a person.

WHAT THE AGENT DOES

Checks the account, applies the policy, completes the action and confirms with the customer.

WATCH OUT

Make escalation easy and tell customers when they are dealing with AI.

ASSISTANT OR AGENT?

Not every job needs an agent

Agents add power and complexity. Often a well-briefed assistant, or a simple automation, is the better answer.

QUESTION	AN ASSISTANT IS ENOUGH WHEN	AN AGENT EARNS ITS PLACE WHEN
Steps	The job is one step: draft, summarise, explain.	The job takes several steps across tools or systems.
Systems	Everything needed can be pasted in.	It needs to look things up or update records itself.
Volume	It happens occasionally, for one person.	It happens often, in a queue, for a team.
Predictability	Every case is different and needs judgement.	Most cases follow a pattern, with clear exceptions.
Checking	A person reads the output anyway.	Results can be checked by rules, tests or sampling.

PRIORITISE

Where to start

Plot candidate agent tasks by value and by how easily a mistake can be caught and undone. Begin top right.

↑ VALUE

DESIGN CAREFULLY

High value, hard to reverse

Payments, customer commitments, production changes. Keep approval steps and go slowly.

START HERE

High value, reversible

Triage, research, data preparation and drafts. Clear owners, easy to check, easy to undo.

AVOID FOR NOW

Low value, hard to reverse

Little to gain and real potential for harm. Revisit when tools and controls mature.

LEARN SAFELY

Low value, reversible

Good sandboxes for teams to build skills and test platforms.

EASE OF CATCHING AND REVERSING MISTAKES →

03

PART 03

Safe by design.

The permissions, approvals, logs and tests that let you give agents real work with confidence.

01 Guardrails that matter

02 Prompt injection and connections

03 Ownership and measurement

GUARDRAILS

Five guardrails for every agent

These apply whatever platform you use. They are the agent equivalent of good access management.

01 **Least privilege**

Grant only the tools and data the task needs. Read-only by default, with write access added deliberately and scoped tightly.

02 **Human approval at the moments that matter**

Payments, deletions, external messages and changes to access go to a person first. Make approval quick so it is not bypassed.

03 **A full audit trail**

Log every tool call, input and result, under an identity that clearly shows it was the agent. You need to know what happened and why.

04 **Test before trust**

Run the agent on realistic cases, including awkward and hostile ones, before it touches live work. Retest after every change.

05 **A clear off switch**

Anyone responsible must be able to pause the agent instantly, and there must be a manual fallback for the work it does.

UNTRUSTED CONTENT

Prompt injection: when the data gives the orders

> If an agent reads content you do not control, assume someone will eventually try to give it instructions.

Agents read things: emails, web pages, documents, tickets, chat messages. The model behind the agent cannot always tell the difference between content it should process and instructions it should follow. That opens a door called **prompt injection**, where text hidden in a document or web page tries to redirect the agent.

A simple example: an agent summarising inbound emails meets a message containing hidden text telling it to forward the mailbox contents to an outside address. If the agent has permission to send email, there is a real risk it tries. AI providers are working hard on defences, but no one currently claims the problem is fully solved.

The most reliable protection is not clever prompting. It is design. Keep agents that read untrusted content away from sensitive actions, or put a human approval step between them. Limit where an agent can send data. Do not combine access to private data, exposure to untrusted content and the ability to communicate externally in one agent unless there is a strong control in the middle.

Testing should include this threat. When you build an agent's test set, add documents and messages that contain misleading or hostile instructions, and check that the agent ignores them or flags them for a person. Repeat those tests whenever the model, the tools or the instructions change, because behaviour that was safe last month is not guaranteed to stay safe.

Connecting agents to tools is getting easier. The **Model Context Protocol (MCP)** is an open standard, originally introduced by Anthropic, that gives AI applications a common way to connect to tools and data sources. It is now supported by many AI tools and platforms, and many software vendors offer MCP connectors for their products.

Standards like MCP are welcome because they reduce custom integration work. But each connector is still a piece of software with access to your systems. Treat it like any other integration: know who built it, review what it can do, approve it centrally and keep a list of what is connected to what. Remove connectors that are no longer needed, and be wary of connectors from unknown sources that ask for far broader access than the job requires.

01

Permissions and least privilege

Most agent incidents are permission problems. Get the access model right and most other risks shrink with it.

IDENTITY

Give each agent its own named account or service identity, never a person's login, so actions are traceable and access can be revoked cleanly.

SCOPE

List the exact tools and actions the task needs. Prefer narrow tools, such as update this field, over broad ones, such as run any command.

LIMITS

Set boundaries in the system, not just in the prompt: value caps, record types, rate limits and approved recipients.

REVIEW

Recertify agent access on the same cycle as staff access, and remove tools that are no longer used.

OWNER

IT or security, with the business process owner

Bring this to the monthly review, with a date for each action.

GOING LIVE

Six steps from idea to trusted agent

A deliberate path that builds evidence at each stage, so autonomy is earned rather than assumed.

01 Name the job and the owner

Write down the task, the expected output and the person accountable for the agent's work.

02 Map tools and permissions

List every system it touches and the minimum access needed. Get security sign-off.

03 Build a test set

Collect realistic cases, including edge cases and deliberate attempts to mislead it.

04 Run in suggest mode

Let it propose actions on live work while people compare its choices with their own.

05 Raise autonomy in steps

Move to approval, then to limits, only when the results justify it.

06 Monitor and review

Check logs, error rates and feedback regularly. Retest after any change to model, tools or instructions.

REMEMBER

Ownership does not transfer to the agent. Every agent needs a named person who answers for what it does.

MEASURING AGENTS

How to tell if an agent is working

Measure the work, not the technology. These measures tell you whether to expand, adjust or retire an agent.

MEASURE	WHAT IT TELLS YOU	HOW TO CAPTURE IT
Completion	How often the agent finishes the task without help.	Agent logs and hand-over counts.
Accuracy	Whether what it did was right.	Regular sampling and reviewer checks.
Time to outcome	Whether work moves faster end to end.	Ticket or case timestamps, before and after.
Escalations	Where it gets stuck or hands over.	Hand-over reasons, reviewed monthly.
Incidents	Whether guardrails are holding.	Blocked actions, near misses and errors logged.
User view	Whether staff and customers trust it.	Short feedback prompts and complaints.

BEFORE YOU START

Questions for your leadership team

Agree the answers before any agent is given access to live systems.

VALUE AND OWNERSHIP

? Which repetitive, multi-step processes cost us the most time today?

? Who will own each agent's results, and its mistakes?

? What evidence would let us give an agent more autonomy?

CONTROL AND RISK

? Which actions must always need a human approval?

? Can we see, in one place, what every agent can access and has done?

? If an agent misbehaved tonight, how quickly could we stop it?

CHECKLIST

Are you ready to let AI act?

Tick what is in place. Gaps are your to-do list before the first agent goes live.

- [ ] **A named owner for each agent**
  Accountable for its work, as for a team member's.

- [ ] **Agents run under their own identities**
  Never under a person's account.

- [ ] **Least-privilege access, reviewed regularly**
  Read-only by default, write access scoped tightly.

- [ ] **Approval steps for high-impact actions**
  Payments, deletions, external messages, access changes.

- [ ] **Every action logged and reviewable**
  With logs kept long enough to investigate issues.

- [ ] **A tested off switch and manual fallback**
  Everyone involved knows how to use it.

- [ ] **A central register of agents and connectors**
  Including what each one can reach.

GLOSSARY

Words worth knowing

Agent

An AI system that uses tools to take a series of actions towards a goal.

Tool

A capability an agent can call, such as search, read a file or update a record.

Autonomy

How much an agent may do without a person approving each step.

Least privilege

Giving an agent only the access its task needs, and no more.

Human in the loop

A person reviews or approves an agent's action before it takes effect.

Audit log

A record of every action an agent took, with inputs and results.

Prompt injection

Hidden or malicious instructions in content an agent reads, aiming to redirect it.

MCP

Model Context Protocol, an open standard for connecting AI applications to tools and data.

Connector

An integration that gives an assistant or agent access to another system.

Runbook

A written procedure for a routine technical task, often a safe starting point for agents.

Sandbox

An isolated environment where an agent can be tested without touching live systems.

TRANSPUTEC

Ready to put agents to work, safely?

We help organisations choose the right first agents, connect them securely to the systems they need and build the controls that keep them trustworthy.

Book an AI agents workshop

transputec.com

ALSO IN THE SERIES

No. 09 · Which AI, for which job?